

NLP Tasks: Question Answering (QA)

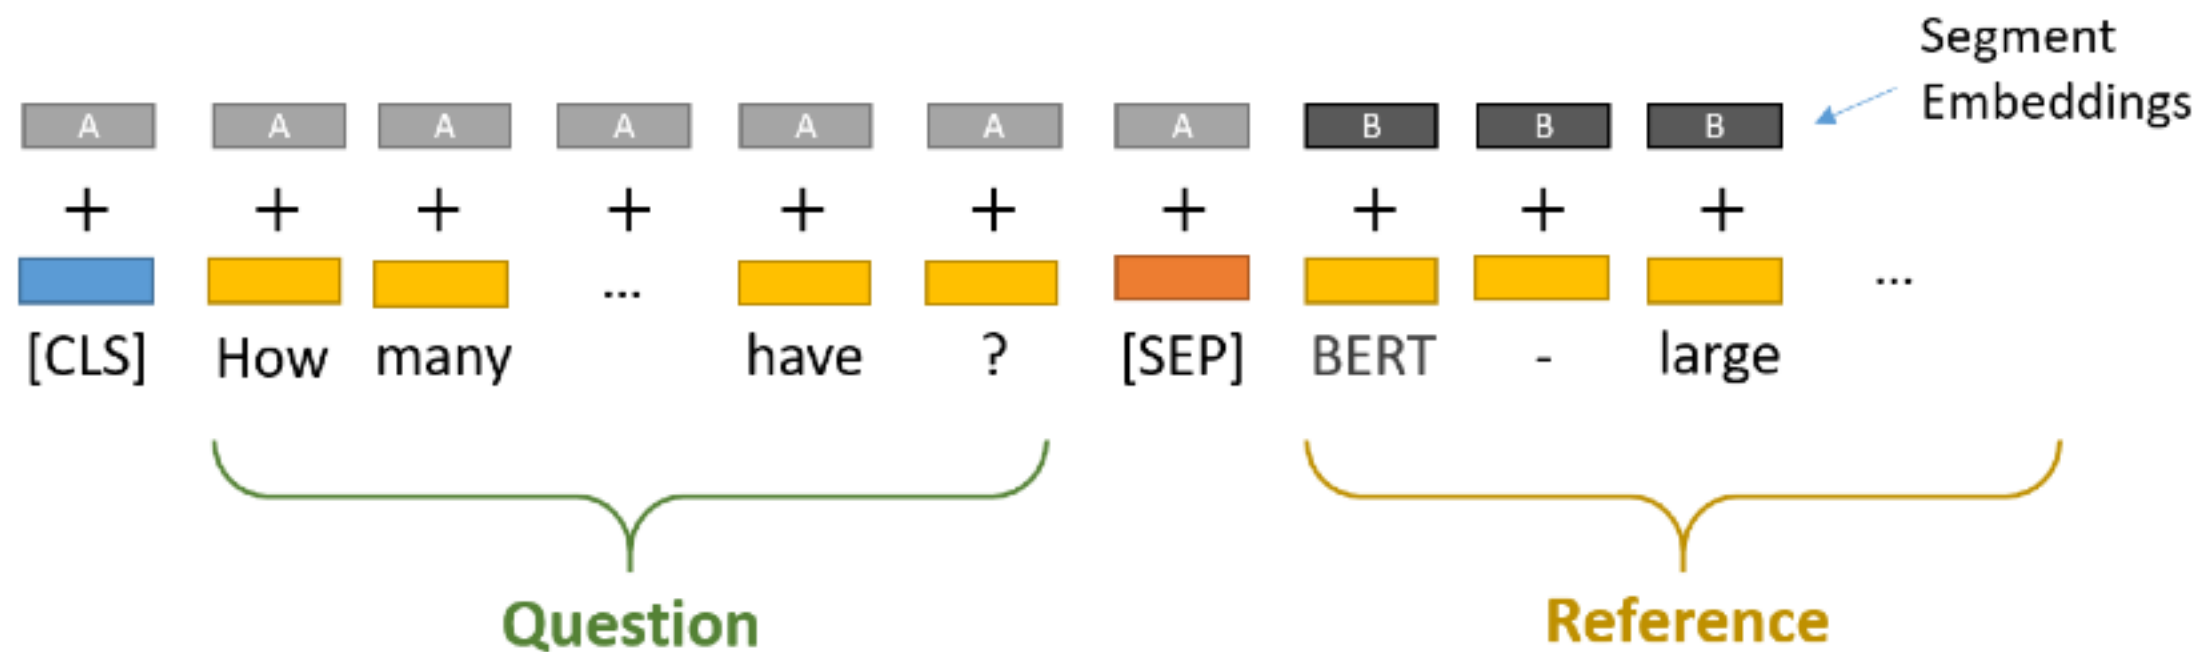


- Question answering is a task by which we automate the selection of a span of text in some corpus that answers some question.
- SQuAD V2.0 is the current benchmark task for evaluating and comparing models.
- Model performance eclipsed human level performance circa 2018.
- Example:

Question Answering (QA) leaderboard



credit:christianclinic



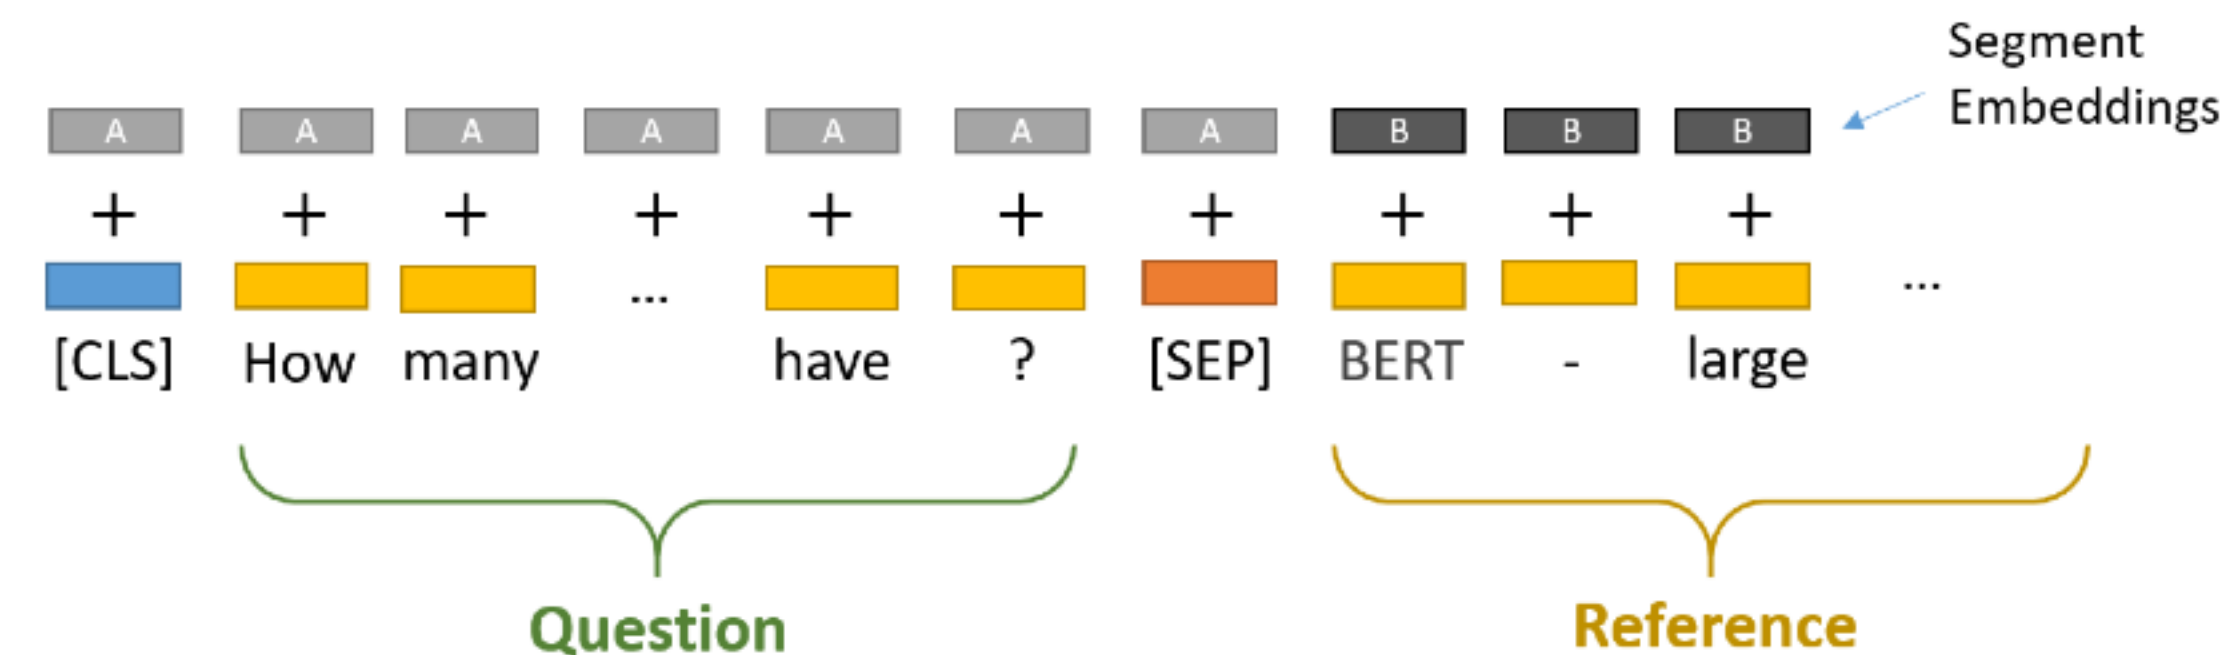
Question: How many parameters does BERT-large have?

Reference Text: BERT-large is really big... it has 24 layers and an embedding size of 1,024, for a total of 340M parameters! Altogether it is 1.34GB, so expect it to take a couple minutes to download to your Colab instance.

NLP Tasks: Question Answering (QA)

- Question answering is a task by which we automate the selection of a span of text in some corpus that answers some question.
- SQuAD V2.0 is the current benchmark task for evaluating and comparing models.
- Model performance eclipsed human level performance circa 2018.

- Example:



Question: How many parameters does BERT-large have?

Reference Text: BERT-large is really big... it has 24 layers and an embedding size of 1,024, for a total of 340M parameters! Altogether it is 1.34GB, so expect it to take a couple minutes to download to your Colab instance.

Question Answering (QA) leaderboard

NLP Tasks: Text summarization

- Text summarization is the task of taking a passage of text as input and generating a summary containing the salient points in the original passage.
- Two approaches to this:
 - Extractive: Returns spans from the original passage
 - Abstractive: Generates a unique text sequence
- Example:

PREMIS:	A soccer game with multiple males playing
HYPOTHESIS:	Some men are playing a sport
LABEL:	Entailment