# Task based assistant core challenges: system scale

Robust language sensing
* Spoken vs typed
* Slang vs formal
* Typos & dialects
* Microphone quality
* Intent, entity, domain rec.
* Implicit entities

Dialogue state management
* Efficient slot filling
* Multi intent requests
* Contextual awareness

**System scale**
* **Multi-domain support**
* **3rd party integrations**
* **Performance**

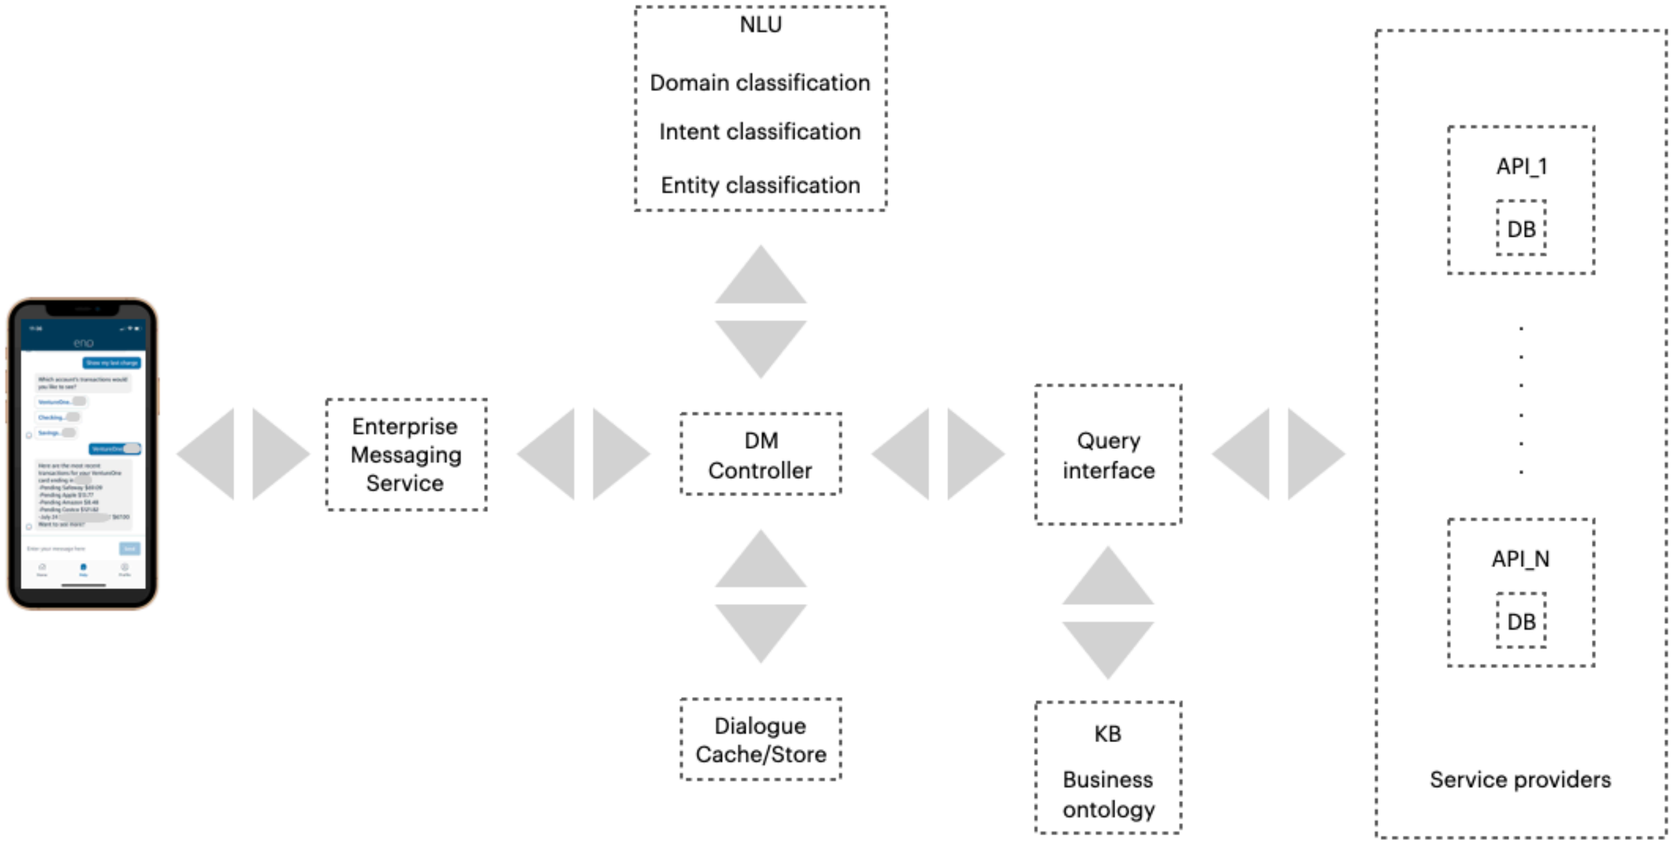


Generic dialogue system architecture

# Scaling dialogue systems involve many considerations

* Understand system level latency profiles
    * Messaging bus
    * Internal APIs that you don't own

* Compute-bound services should be isolated
    * Dynamic batch processing

* Unit testing, including 3rd party tools

* Understand / measure the latency-throughput curve

* Cost
    * Cost-performance-headache landscape
    * Accelerators cost (a lot) more!
    * Use hardware designed for inference, not training.

* Set SLAs for 3rd party APIs
    * Define pXX latency bounds (and enforce them)